

Package ‘OptHoldoutSize’

December 20, 2021

Title Estimation of Optimal Size for a Holdout Set for Updating a Predictive Score

Version 0.0.0.9000

Author James Liley, Sam Emerson, Sami Haidar-Wehbe, and Louis Aslett

Maintainer James Liley <james.liley@durham.ac.uk>

Description Predictive scores must be updated with care, because actions taken on the basis of existing risk scores causes bias in risk estimates from the updated score. A holdout set is a straightforward way to manage this problem: a proportion of the population is 'held-out' from computation of the previous risk score. This package provides tools to estimate a size for this holdout set and associated errors. Comprehensive vignettes are included.

License GPL (>=3)

Encoding UTF-8

LazyData true

Depends R (>= 3.5.0), matrixStats, mnormt, mvtnorm, ranger, mle.tools

Roxygen list(markdown = TRUE)

RoxygenNote 7.1.2

Suggests knitr, rmarkdown

VignetteBuilder knitr

NeedsCompilation no

R topics documented:

| | |
|-----------------------------------|----|
| add_aspre_interactions | 2 |
| aspre | 3 |
| aspre_emulation | 4 |
| aspre_k2 | 4 |
| aspre_parametric | 5 |
| ci_cover_a_yn | 5 |
| ci_cover_e_yn | 6 |
| ci_ohs | 6 |
| cov_fn | 9 |
| data_example_simulation | 10 |
| data_nextpoint_em | 10 |
| data_nextpoint_par | 11 |
| error_ohs_emulation | 11 |

| | |
|--|-----------|
| exp_imp_fn | 13 |
| gen_base_coefs | 15 |
| gen_preds | 16 |
| gen_resp | 16 |
| grad_nstar_powerlaw | 17 |
| logistic | 18 |
| logit | 18 |
| model_predict | 19 |
| model_train | 20 |
| mu_fn | 21 |
| next_n | 23 |
| ohs_array | 25 |
| ohs_resample | 25 |
| optimal_holdout_size | 26 |
| optimal_holdout_size_emulation | 27 |
| oracle_pred | 28 |
| params_aspre | 29 |
| plot.optholdoutsizes | 29 |
| plot.optholdoutsizes_emul | 30 |
| powersolve | 31 |
| powersolve_general | 32 |
| powersolve_se | 32 |
| psi_fn | 34 |
| sens10 | 35 |
| sim_random_aspre | 36 |
| split_data | 37 |
| Index | 38 |

add_aspre_interactions

Add interaction terms corresponding to ASPRE model

Description

Add various interaction terms to X. Interaction terms correspond to those in ASPRE.

Usage

```
add_aspre_interactions(X)
```

Arguments

X data frame

Value

New data frame containing interaction terms.

Examples

```
# Load ASPRE related data
data(params_aspre)

X=sim_random_aspre(1000,params_aspre)
Xnew=add_aspre_interactions(X)

print(colnames(X))
print(colnames(Xnew))
```

| | |
|-------|-----------------------------|
| aspre | <i>Computes ASPRE score</i> |
|-------|-----------------------------|

Description

Computes ASPRE model prediction on a matrix X of covariates

Full ASPRE model from https://www.nejm.org/doi/suppl/10.1056/NEJMoa1704559/suppl_file/nejmoa1704559_append

Model is to predict gestational age at PE; that is, a higher score indicates a lower PE risk, so coefficients are negated for model to predict PE risk.

Usage

```
aspre(X)
```

Arguments

X matrix, assumed to be output of `sim_random_aspre` with parameter `params=params_aspre` and transformed using `add_aspre_interactions`

Value

vector of scores.

Examples

```
# Load ASPRE related data
data(params_aspre)

X=sim_random_aspre(1000,params_aspre)
Xnew=add_aspre_interactions(X)

aspre_score=aspre(Xnew)

plot(density(aspre_score))
```

| | |
|-----------------|---|
| aspre_emulation | <i>Emulation-based OHS estimation for ASPRE</i> |
|-----------------|---|

Description

This object contains data relating to emulation-based OHS estimation for the ASPRE model. For generation, see hidden code in vignette, or in pipeline at https://github.com/jamesliley/OptHoldoutSize_pipelines

Usage

```
aspre_emulation
```

Format

An object of class list of length 4.

| | |
|----------|---|
| aspre_k2 | <i>Cost estimating function in ASPRE simulation</i> |
|----------|---|

Description

Estimate cost at a given holdout set size in ASPRE model

Usage

```
aspre_k2(
  n,
  X,
  PRE,
  seed = NULL,
  pi_PRE = 1426/58974,
  pi_intervention = 0.1,
  alpha = 0.37
)
```

Arguments

| | |
|-----------------|--|
| n | Holdout set size at which to estimate k_2 (cost) |
| X | Matrix of predictors |
| PRE | Vector indicating PRE incidence |
| seed | Random seed; set before starting or set to NULL |
| pi_PRE | Population prevalence of PRE if not prophylactically treated. Defaults to empirical value 1426/58974 |
| pi_intervention | Proportion of the population on which an intervention will be made. Defaults to 0.1 |
| alpha | Reduction in PRE risk with intervention. Defaults to empirical value 0.37 |

Value

Estimated cost

Examples

```
# Simulate
set.seed(32142)

N=1000; p=15
X=matrix(rnorm(N*p),N,p); PRE=rbinom(N,1,prob=logit(X%% rnorm(p)))
aspre_k2(1000,X,PRE)
```

| | |
|------------------|--|
| aspre_parametric | <i>Parametric-based OHS estimation for ASPRE</i> |
|------------------|--|

Description

This object contains data relating to parametric-based OHS estimation for the ASPRE model. For generation, see hidden code in vignette, or in pipeline at https://github.com/jamesliley/OptHoldoutSize_pipelines

Usage

```
aspre_parametric
```

Format

An object of class `list` of length 4.

| | |
|---------------|---|
| ci_cover_a_yn | <i>Data for example on asymptotic confidence interval for CI.</i> |
|---------------|---|

Description

Data for example for asymptotic confidence interval for CI. For generation, see example.

Usage

```
ci_cover_a_yn
```

Format

An object of class `matrix` (inherits from `array`) with 11 rows and 5000 columns.

| | |
|---------------|--|
| ci_cover_e_yn | <i>Data for example on empirical confidence interval for CI.</i> |
|---------------|--|

Description

Data for example for empirical confidence interval for CI. For generation, see example.

Usage

```
ci_cover_e_yn
```

Format

An object of class `matrix` (inherits from `array`) with 11 rows and 5000 columns.

| | |
|--------|---|
| ci_ohs | <i>Confidence interval for optimal holdout size, when estimated using parametric method</i> |
|--------|---|

Description

Compute confidence interval for optimal holdout size given either a standard error covariance matrix or a set of `n_e` estimates of parameters.

This can be done either asymptotically, using a method analogous to the Fisher information matrix, or empirically (using bootstrap resampling)

If `sigma` (covariance matrix) is specified and `method='bootstrap'`, a confidence interval is generated assuming a Gaussian distribution of $(N, k1, \theta)$. To estimate a confidence interval assuming a non-Gaussian distribution, simulate values under the requisite distribution and use then as parameters $N, k1, \theta$, with `sigma` set to `NULL`.

Usage

```
ci_ohs(
  N,
  k1,
  theta,
  alpha = 0.05,
  k2 = powerlaw,
  grad_nstar = NULL,
  sigma = NULL,
  n_boot = 1000,
  seed = NULL,
  mode = "empirical",
  ...
)
```

Arguments

| | |
|------------|--|
| N | Vector of estimates of total number of samples on which the predictive score will be used/fitted, or single estimate |
| k1 | Vector of estimates of cost value in the absence of a predictive score, or single number |
| theta | Matrix of estimates of parameters for function $k_2(n)$ governing expected cost to an individual sample given a predictive score fitted to n samples. Can be a matrix of dimension $n \times n_{\text{par}}$, where n_{par} is the number of parameters of k_2 . |
| alpha | Construct $1-\alpha$ confidence interval. Defaults to 0.05 |
| k2 | Function governing expected cost to an individual sample given a predictive score fitted to n samples. Must take two arguments: n (number of samples) and θ (parameters). Defaults to a power-law form $k_2(n, c(a, b, c)) = a n^{-(b)} + c$. |
| grad_nstar | Function giving partial derivatives of optimal holdout set, taking three arguments: N , k_1 , and θ . Only used for asymptotic confidence intervals. If NULL, estimated empirically |
| sigma | Standard error covariance matrix for (N, k_1, θ) , in that order. If NULL, will derive as sample covariance matrix of parameters. Must be of the correct size and positive definite. |
| n_boot | Number of bootstrap resamples for empirical estimate. |
| seed | Random seed for bootstrap resamples. Defaults to NULL. |
| mode | One of 'asymptotic' or 'empirical'. Defaults to 'empirical' |
| ... | Passed to function optimize |

Value

A vector of length two containing lower and upper limits of confidence interval.

Examples

```
## We will assume that our observations of N, k1, and theta=(a,b,c) are distributed with mean mu_par and variance
mu_par=c(N=10000,k1=0.35,A=3,B=1.5,C=0.1)
sigma_par=cbind(
  c(100^2, 1, 0, 0, 0),
  c( 1, 0.07^2, 0, 0, 0),
  c( 0, 0, 0.5^2, 0.05, -0.001),
  c( 0, 0, 0.05, 0.4^2, -0.002),
  c( 0, 0, -0.001, -0.002, 0.02^2)
)

# Firstly, we make 500 observations
par_obs=rmnorm(500,mean=mu_par,varcov=sigma_par)

# Optimal holdout size and asymptotic and empirical confidence intervals
ohs=optimal_holdout_size(N=mean(par_obs[,1]),k1=mean(par_obs[,2]),theta=colMeans(par_obs[,3:5]))$size
ci_a=ci_ohs(N=par_obs[,1],k1=par_obs[,2],theta=par_obs[,3:5],alpha=0.05,seed=12345,mode="asymptotic")
ci_e=ci_ohs(N=par_obs[,1],k1=par_obs[,2],theta=par_obs[,3:5],alpha=0.05,seed=12345,mode="empirical")

# Assess cover at various n_e
n_e_values=c(20,30,50,100,150,200,300,500,750,1000,1500)
ntrial=5000
```

```

alpha_trial=0.1 # use 90% confidence intervals
nstar_true=optimal_holdout_size(N=mu_par[1],k1=mu_par[2],theta=mu_par[3:5])$size

## The matrices indicating cover take are included in this package but take around 30 minutes to generate. They a
data(ci_cover_a_yn)
data(ci_cover_e_yn)

if (!exists("ci_cover_a_yn")) {
  ci_cover_a_yn=matrix(NA,length(n_e_values),ntrial) # Entry [i,j] is 1 if ith asymptotic CI for jth value of n_e
  ci_cover_e_yn=matrix(NA,length(n_e_values),ntrial) # Entry [i,j] is 1 if ith empirical CI for jth value of n_e

  for (i in 1:length(n_e_values)) {
    n_e=n_e_values[i]
    for (j in 1:ntrial) {
      # Set seed
      set.seed(j*ntrial + i + 12345)

      # Make n_e observations
      par_obs=rnorm(n_e,mean=mu_par,varcov=sigma_par)
      ci_a=ci_ohs(N=par_obs[,1],k1=par_obs[,2],theta=par_obs[,3:5],alpha=alpha_trial,mode="asymptotic")
      ci_e=ci_ohs(N=par_obs[,1],k1=par_obs[,2],theta=par_obs[,3:5],alpha=alpha_trial,mode="empirical",n_boot=

      if (nstar_true>ci_a[1] & nstar_true<ci_a[2]) ci_cover_a_yn[i,j]=1 else ci_cover_a_yn[i,j]=0
      if (nstar_true>ci_e[1] & nstar_true<ci_e[2]) ci_cover_e_yn[i,j]=1 else ci_cover_e_yn[i,j]=0
    }
    print(paste0("Completed for n_e = ",n_e))
  }
}

# Cover at each n_e value and standard error
cover_a=rowMeans(ci_cover_a_yn)
cover_e=rowMeans(ci_cover_e_yn)
zse_a=2*sqrt(cover_a*(1-cover_a)/ntrial)
zse_e=2*sqrt(cover_e*(1-cover_e)/ntrial)

# Draw plot. Convergence to 1-alpha cover is evident. Cover is not far from alpha even at small n_e.

plot(0,type="n",xlim=range(n_e_values),ylim=c(0.7,1),xlab=expression("n"[e]),ylab="Cover")

# Asymptotic cover and 2*SE pointwise envelope
polygon(c(n_e_values,rev(n_e_values)),c(cover_a+zse_a,rev(cover_a-zse_a)),
  col=rgb(1,1,1,alpha=0.3),border=NA)
lines(n_e_values,cover_a,col="black")

# Empirical cover and 2*SE pointwise envelope
polygon(c(n_e_values,rev(n_e_values)),c(cover_e+zse_e,rev(cover_e-zse_e)),
  col=rgb(0,0,1,alpha=0.3),border=NA)
lines(n_e_values,cover_e,col="blue")

abline(h=1-alpha_trial,col="red")
legend("bottomright",c("Asym.", "Emp.", expression(paste("1-",alpha))),lty=1,col=c("black", "blue", "red"))

```


cov_fn

*Covariance function for Gaussian process***Description**

Radial kernel covariance function for Gaussian process.

Used for a Gaussian process $GP(m, k(\cdot, \cdot))$, an instance X of which has covariance $k(n, n')$ between $X(n)$ and $X(n')$.

Covariance function is parametrised by `var_u` (general variance) and `k_width` (kernel width)

Usage

```
cov_fn(n, ndash, var_u, k_width)
```

Arguments

| | |
|----------------------|--|
| <code>n</code> | Argument 1: kernel is a function of <code>ndash-n</code> |
| <code>ndash</code> | Argument 2: kernel is a function of <code>ndash-n</code> |
| <code>var_u</code> | Global variance |
| <code>k_width</code> | Kernel width |

Value

Covariance value

Examples

```
##' # We will sample from Gaussian processes  $GP(0, k(\cdot, \cdot) = \text{cov\_fn}(\cdot, \cdot; \text{var\_u}, \text{theta}))$  at these values of n
nvals=seq(1, 300, length=100)

# We will consider two theta values
kw1=10; kw2=30

# We will consider two var_u values
var1=1; var2=10

# Covariance matrices
cov11=outer(nvals, nvals, function(n, ndash) cov_fn(n, ndash, var_u=var1, k_width=kw1))
cov12=outer(nvals, nvals, function(n, ndash) cov_fn(n, ndash, var_u=var1, k_width=kw2))
cov21=outer(nvals, nvals, function(n, ndash) cov_fn(n, ndash, var_u=var2, k_width=kw1))
cov22=outer(nvals, nvals, function(n, ndash) cov_fn(n, ndash, var_u=var2, k_width=kw2))

# Dampen slightly to ensure positive definiteness
damp=1e-5
cov11=(1-damp)*(1-diag(length(nvals)))*cov11 + diag(length(nvals))*cov11
cov12=(1-damp)*(1-diag(length(nvals)))*cov12 + diag(length(nvals))*cov12
cov21=(1-damp)*(1-diag(length(nvals)))*cov21 + diag(length(nvals))*cov21
cov22=(1-damp)*(1-diag(length(nvals)))*cov22 + diag(length(nvals))*cov22

# Sample
set.seed(35243)
y11=rmnorm(1, mean=0, varcov=cov11)
```

```
y12=rmnorm(1,mean=0,varcov=cov12)
y21=rmnorm(1,mean=0,varcov=cov21)
y22=rmnorm(1,mean=0,varcov=cov22)

# Plot
rr=max(abs(c(y11,y12,y21,y22)))
plot(0,xlim=range(nvals),ylim=c(-rr,rr+10),xlab="n",ylab=expression("GP(0,cov_fn(.,.;var_u,theta)"))
lines(nvals,y11,lty=1,col="black")
lines(nvals,y12,lty=2,col="black")
lines(nvals,y21,lty=1,col="red")
lines(nvals,y22,lty=2,col="red")
legend("topright",c("k_width=10, var_u=1", "k_width=30, var_u=1", "k_width=10, var_u=10", "k_width=30, var_u=10"),
      lty=c(1,2,1,2),col=c("black","black","red","red"))
```

data_example_simulation

Data for vignette showing general example

Description

Data for general vignette. For generation, see hidden code in vignette, or in pipeline at <https://github.com/jamesliley/OptH>

Usage

```
data_example_simulation
```

Format

An object of class list of length 4.

data_nextpoint_em

Data for 'next point' demonstration vignette on algorithm comparison using emulation algorithm

Description

Data containing 'next point selected' information for emulation algorithm in vignette comparing emulation and parametric algorithms. For generation, see hidden code in vignette, or in pipeline at https://github.com/jamesliley/OptHoldoutSize_pipelines

Usage

```
data_nextpoint_em
```

Format

An object of class list of length 6.

| | |
|--------------------|--|
| data_nextpoint_par | <i>Data for 'next point' demonstration vignette on algorithm comparison using parametric algorithm</i> |
|--------------------|--|

Description

Data containing 'next point selected' information for parametric algorithm in vignette comparing emulation and parametric algorithms. For generation, see hidden code in vignette, or in pipeline at https://github.com/jamesliley/OptHoldoutSize_pipelines

Usage

```
data_nextpoint_par
```

Format

An object of class `list` of length 6.

| | |
|---------------------|---|
| error_ohs_emulation | <i>Measure of error for emulation-based OHS emulation</i> |
|---------------------|---|

Description

Measure of error for semiparametric (emulation) based estimation of optimal holdout set sizes.

Returns a set of values of n for which a $1-\alpha$ credible interval for cost at includes a lower value than the cost at the estimated optimal holdout size.

This is not a confidence interval, credible interval or credible set for the OHS, and is prone to misinterpretation.

Usage

```
error_ohs_emulation(
  nset,
  d,
  var_w,
  N,
  k1,
  alpha = 0.1,
  var_u = 1e+07,
  k_width = 5000,
  mean_fn = powerlaw,
  theta = powersolve(nset, d, y_var = var_w)$par,
  npoll = 1000
)
```

Arguments

| | |
|---------|--|
| nset | Training set sizes for which a loss has been evaluated |
| d | Loss at training set sizes nset |
| var_w | Variance of error in loss estimate at each training set size. |
| N | Total number of samples on which the model will be fitted/used |
| k1 | Mean loss per sample with no predictive score in place |
| alpha | Use 1-alpha credible interval. Defaults to 0.1. |
| var_u | Marginal variance for Gaussian process kernel. Defaults to 1e7 |
| k_width | Kernel width for Gaussian process kernel. Defaults to 5000 |
| mean_fn | Functional form governing expected loss per sample given sample size. Should take two parameters: n (sample size) and theta (parameters). Defaults to function powerlaw. |
| theta | Current estimates of parameter values for mean_fn. Defaults to the MLE power-law solution corresponding to n,d, and var_w. |
| npoll | Check npoll equally spaced values between 1 and N for minimum. If NULL, check all values (this can be slow). Defaults to 1000 |

Value

Vector of values n for which 1-alpha credible interval for cost $l(n)$ at n contains mean posterior loss at estimated optimal holdout size.

Examples

```
# Set seed
set.seed(57365)

# Parameters
N=100000;
k1=0.3
A=8000; B=1.5; C=0.15; theta=c(A,B,C)

# True mean function
k2_true=function(n) powerlaw(n,theta)

# True OHS
nx=1:N
ohs_true=nx[which.min(k1*nx + k2_true(nx)*(N-nx))]

# Values of n for which cost has been estimated
np=50 # this many points
nset=round(runif(np,1,N))
var_w=runif(np,0.001,0.0015)
d=rnorm(np,mean=k2_true(nset),sd=sqrt(var_w))

# Compute OHS
res1=optimal_holdout_size_emulation(nset,d,var_w,N,k1)

# Error estimates
ex=error_ohs_emulation(nset,d,var_w,N,k1)

# Plot
```

```

plot(res1)

# Add error
abline(v=ohs_true)
abline(v=ex,col=rgb(1,0,0,alpha=0.2))

# Show justification for error
n=seq(1,N,length=1000)
mu=mu_fn(n,nset,d,var_w,N,k1); psi=pmax(0,psi_fn(n, nset, var_w, N)); Z=-qnorm(0.1/2)
lines(n,mu - Z*sqrt(psi),lty=2,lwd=2)
legend("topright",
      c("Err. region",expression(paste(mu(n)- "z"[alpha/2]*sqrt(psi(n))))),
      pch=c(16,NA),lty=c(NA,2),lwd=c(NA,2),col=c("pink","black"),bty="n")

```

| | |
|------------|-----------------------------|
| exp_imp_fn | <i>Expected improvement</i> |
|------------|-----------------------------|

Description

Expected improvement

Essentially chooses the next point to add to n , called n^* , in order to minimise the expectation of $\text{loss}(n^*)$.

Usage

```

exp_imp_fn(
  n,
  nset,
  d,
  var_w,
  N,
  k1,
  var_u = 1e+07,
  k_width = 5000,
  mean_fn = powerlaw,
  theta = powersolve(nset, d, y_var = var_w)$par
)

```

Arguments

| | |
|----------------------|--|
| <code>n</code> | Set of training set sizes to evaluate |
| <code>nset</code> | Training set sizes for which a loss has been evaluated |
| <code>d</code> | Loss at training set sizes <code>nset</code> |
| <code>var_w</code> | Variance of error in loss estimate at each training set size. |
| <code>N</code> | Total number of samples on which the model will be fitted/used |
| <code>k1</code> | Mean loss per sample with no predictive score in place |
| <code>var_u</code> | Marginal variance for Gaussian process kernel. Defaults to $1e7$ |
| <code>k_width</code> | Kernel width for Gaussian process kernel. Defaults to 5000 |

| | |
|---------|--|
| mean_fn | Functional form governing expected loss per sample given sample size. Should take two parameters: n (sample size) and theta (parameters). Defaults to function powerlaw. |
| theta | Current estimates of parameter values for mean_fn. Defaults to the MLE power-law solution corresponding to n,d, and var_w. |

Value

Value of expected improvement at values n

Examples

```
# Set seed.
set.seed(24015)

# Kernel width and Gaussian process variance
kw0=5000
vu0=1e7

# Include legend on plots or not; inclusion can obscure plot elements on small figures
inc_legend=FALSE

# Suppose we have population size and cost-per-sample without a risk score as follows
N=100000
k1=0.4

# Suppose that true values of a,b,c are given by
theta_true=c(10000,1.2,0.2)
theta_lower=c(1,0.5,0.1) # lower bounds for estimating theta
theta_upper=c(20000,2,0.5) # upper bounds for estimating theta

# We start with five random holdout set sizes (nset0),
# with corresponding cost-per-individual estimates d0 derived
# with various errors var_w0
nstart=4
vwmin=0.001; vwmax=0.005
nset0=round(runif(nstart,1000,N/2))
var_w0=runif(nstart,vwmin,vwmax)
d0=rnorm(nstart,mean=powerlaw(nset0,theta_true),sd=sqrt(var_w0))

# We estimate theta from these three points
theta0=powersolve(nset0,d0,y_var=var_w0,lower=theta_lower,upper=theta_upper,init=theta_true)$par

# We will estimate the posterior at these values of n
n=seq(1000,N,length=1000)

# Mean and variance
p_mu=mu_fn(n,nset=nset0,d=d0,var_w = var_w0, N=N,k1=k1,theta=theta0,k_width=kw0,var_u=vu0)
p_var=psi_fn(n,nset=nset0,N=N,var_w = var_w0,k_width=kw0,var_u=vu0)

# Plot
yrange=c(-30000,100000)
plot(0,xlim=range(n),ylim=yrange,type="n",
     xlab="Training/holdout set size",
```

```

      ylab="Total cost (= num. cases)")
      lines(n,p_mu,col="blue")
      lines(n,p_mu - 3*sqrt(p_var),col="red")
      lines(n,p_mu + 3*sqrt(p_var),col="red")
      points(nset0,k1*nset0 + d0*(N-nset0),pch=16,col="purple")
      lines(n,k1*n + powerlaw(n,theta0)*(N-n),lty=2)
      lines(n,k1*n + powerlaw(n,theta_true)*(N-n),lty=3,lwd=3)
      if (inc_legend) {
        legend("topright",
              c(expression(mu(n)),
                expression(mu(n) %+-% 3*sqrt(psi(n))),
                "prior(n)",
                "True",
                "d"),
              lty=c(1,1,2,3,NA),lwd=c(1,1,1,3,NA),pch=c(NA,NA,NA,NA,16),pt.cex=c(NA,NA,NA,NA,1),
              col=c("blue","red","black","purple"),bg="white")
      }

      ## Add line corresponding to recommended new point
      exp_imp_em <- exp_imp_fn(n,nset=nset0,d=d0,var_w = var_w0, N=N,k1=k1,theta=theta0,k_width=kw0,var_u=vu0)
      abline(v=n[which.max(exp_imp_em)])

```

gen_base_coefs

Coefficients for imperfect risk score

Description

Generate coefficients corresponding to an imperfect risk score, interpretable as 'baseline' behaviour in the absence of a risk score

Usage

```
gen_base_coefs(coefs, noise = TRUE, num_vars = 2, max_base_powers = 1)
```

Arguments

| | |
|-----------------|--|
| coefs | Original coefficients |
| noise | Set to TRUE to add Gaussian noise to coefficients |
| num_vars | Number of variables at hand for baseline calculation |
| max_base_powers | If >1, return a matrix of coefficients, with successively more noise |

Value

Vector of coefficients

Examples

```
# See examples for model_predict
```

| | |
|-----------|---|
| gen_preds | <i>Generate matrix of random observations</i> |
|-----------|---|

Description

Generate matrix of random observations. Observations are unit Gaussian-distributed.

Usage

```
gen_preds(nobs, npreds, ninters = 0)
```

Arguments

| | |
|----------|---|
| nobs | Number of observations (samples) |
| npreds | Number of predictors |
| niinters | Number of interaction terms, default 0. Up to npreds*(npreds-1)/2 |

Value

Data frame of observations

Examples

```
# See examples for model_predict
```

| | |
|----------|--------------------------|
| gen_resp | <i>Generate response</i> |
|----------|--------------------------|

Description

Generate random outcome (response) according to a ground-truth logistic model

Usage

```
gen_resp(X, coefs = NA, coefs_sd = 1, retprobs = FALSE)
```

Arguments

| | |
|----------|---|
| X | Matrix of observations |
| coefs | Vector of coefficients for logistic model. If NA, random coefficients are generated. Defaults to NA |
| coefs_sd | If random coefficients are generated, use this SD (mean 0) |
| retprobs | If TRUE, return class probability; otherwise, return classes. Defaults to FALSE |

Value

Vector of length as first dimension of dim(X) with outcome classes (if retprobs==FALSE) or outcome probabilities (if retprobs==TRUE)

Examples

```
# See examples for model_predict
```

| | |
|---------------------|---|
| grad_nstar_powerlaw | <i>Gradient of optimal holdout size (power law)</i> |
|---------------------|---|

Description

Compute gradient of optimal holdout size assuming a power-law form of k_2

Assumes cost function is $l(n; k_1, N, \theta) = k_1 n + k_2(n; \theta) (N - n)$ with $k_2(n; \theta) = k_2(n; a, b, c) = a n^b (-b) + c$

Usage

```
grad_nstar_powerlaw(N, k1, theta)
```

Arguments

| | |
|-------|---|
| N | Total number of samples on which the predictive score will be used/fitted. Can be a vector. |
| k1 | Cost value in the absence of a predictive score. Can be a vector. |
| theta | Parameters for function $k_2(n)$ governing expected cost to an individual sample given a predictive score fitted to n samples. Can be a matrix of dimension $n \times n_{\text{par}}$, where n_{par} is the number of parameters of k_2 . |

Value

List/data frame of dimension (number of evaluations) \times 5 containing partial derivatives of $nstar$ (optimal holdout size) with respect to N , k_1 , a , b , c respectively.

Examples

```
# Evaluate optimal holdout set size for a range of values of k1, and compute derivative
N=10000;
k1=seq(0.1,0.5,length=20)
A=3; B=1.5; C=0.15; theta=c(A,B,C)

nstar=optimal_holdout_size(N,k1,theta)
grad_nstar=grad_nstar_powerlaw(N,k1,theta)

plot(0,type="n",ylim=c(-2000,500),xlim=range(k1),xlab=expression("k"[1]),ylab="Optimal holdout set size")
lines(nstar$k1,nstar$size,col="black")
lines(nstar$k1,grad_nstar[,2],col="red")
legend("bottomright",c(expression("n"["*"]),expression(paste(partialdiff[k1],"n"["*"]))),
      col=c("black","red"),lty=1)
```

`logistic`*Logistic*

Description

Logistic function: $-\log((1/x)-1)$

Usage

```
logistic(x)
```

Arguments

`x` argument

Value

value of `logit(x)`; na if `x` is outside (0,1)

Examples

```
# Plot
x=seq(0,1,length=100)
plot(x,logistic(x),type="l")

# Logit and logistic are inverses
x=seq(-5,5,length=1000)
plot(x,logistic(logit(x)),type="l")
```

`logit`*Logit*

Description

Logit function: $1/(1+\exp(-x))$

Usage

```
logit(x)
```

Arguments

`x` argument

Value

value of `logit(x)`

Examples

```
# Plot
x=seq(-5,5,length=1000)
plot(x,logit(x),type="l")
```

| | |
|---------------|-------------------------|
| model_predict | <i>Make predictions</i> |
|---------------|-------------------------|

Description

Make predictions according to a given model

Usage

```
model_predict(  
  data_test,  
  trained_model,  
  return_type,  
  threshold = NULL,  
  model_family = NULL,  
  ...  
)
```

Arguments

| | |
|---------------|--|
| data_test | Data for which predictions are to be computed |
| trained_model | Model for which predictions are to be made |
| return_type | ?? |
| threshold | ?? |
| model_family | ?? |
| ... | Passed to function <code>predict.glm()</code> or <code>predict.ranger()</code> |

Value

Vector of predictions

Examples

```
## Set seed for reproducibility  
seed=1234  
set.seed(seed)  
  
# Initialisation of patient data  
n_iter <- 500           # Number of point estimates to be calculated  
nobs <- 5000            # Number of observations, i.e patients  
npreds <- 7             # Number of predictors  
  
# Model family  
family="log_reg"  
  
# Baseline behaviour is an oracle Bayes-optimal predictor on only one variable  
max_base_powers <- 1  
base_vars=1  
  
# Check the following holdout size fractions  
frac_ho = 0.1
```

```

# Set ground truth coefficients, and the accuracy at baseline
coefs_general <- rnorm(npreds,sd=1/sqrt(npreds))
coefs_base <- gen_base_coefs(coefs_general, max_base_powers = max_base_powers)

# Generate dataset
X <- gen_preds(nobs, npreds)

# Generate labels
newdata <- gen_resp(X, coefs = coefs_general)
Y <- newdata$classes

# Combined dataset
pat_data <- cbind(X, Y)
pat_data$Y = factor(pat_data$Y)

# For each holdout size, split data into intervention and holdout set
mask <- split_data(pat_data, frac_ho)
data_interv <- pat_data[!mask,]
data_hold <- pat_data[mask,]

# Train model
trained_model <- model_train(data_hold, model_family = family)
thresh <- 0.5

# Make predictions
class_pred <- model_predict(data_interv, trained_model,
                           return_type = "class",
                           threshold = 0.5, model_family = family)

# Simulate baseline predictions
base_pred <- oracle_pred(data_hold,coefs_base[base_vars, ], num_vars = base_vars)

# Contingency table for model-based predictor (on intervention set)
print(table(class_pred,data_interv$Y))

# Contingency table for model-based predictor (on holdout set)
print(table(base_pred,data_hold$Y))

```

model_train

Train model (wrapper)

Description

Train model using either a GLM or a random forest

Usage

```
model_train(train_data, model_family = "log_reg", ...)
```

Arguments

| | |
|--------------|---|
| train_data | Data to use for training; assumed to have one binary column called Y |
| model_family | Either 'log_reg' for logistic regression or 'rand_forest' for random forest |
| ... | Passed to function glm() or ranger() |

Value

Fitted model of type GLM or Ranger

Examples

```
# See examples for model_predict
```

| | |
|-------|---------------------------|
| mu_fn | <i>Power law function</i> |
|-------|---------------------------|

Description

Power law function for modelling learning curve (taken to mean change in expected loss per sample with training set size)

Recommended in [review of learning curve forms](#)

If $\theta = c(a, b, c)$ then models as $a n^{(-b)} + c$. Note b is negated.

Note that $\text{powerlaw}(n, c(a, b, c))$ has limit c as n tends to infinity, if $a, b > 0$

Posterior mean for emulator given points n.

Usage

```
powerlaw(n, theta)
```

```
mu_fn(
  n,
  nset,
  d,
  var_w,
  N,
  k1,
  var_u = 1e+07,
  k_width = 5000,
  mean_fn = powerlaw,
  theta = powersolve(nset, d, y_var = var_w)$par
)
```

Arguments

| | |
|-------|--|
| n | Set of training set sizes to evaluate |
| theta | Current estimates of parameter values for mean_fn. Defaults to the MLE power-law solution corresponding to n,d, and var_w. |
| nset | Training set sizes for which a loss has been evaluated |

| | |
|---------|--|
| d | Loss at training set sizes nset |
| var_w | Variance of error in loss estimate at each training set size. |
| N | Total number of samples on which the model will be fitted/used |
| k1 | Mean loss per sample with no predictive score in place |
| var_u | Marginal variance for Gaussian process kernel. Defaults to 1e7 |
| k_width | Kernel width for Gaussian process kernel. Defaults to 5000 |
| mean_fn | Functional form governing expected loss per sample given sample size. Should take two parameters: n (sample size) and theta (parameters). Defaults to function powerlaw. |

Value

Vector of values of same length as n

Vector Mu of same length of n where $\text{Mu}_i = \text{mean}(\text{posterior}(\text{cost}(n_i)))$

Examples

```
ncheck=seq(1000,10000)
plot(ncheck, powerlaw(ncheck, c(5e3,1.2,0.3)),type="l",xlab="n",ylab="powerlaw(n)")

# Suppose we have population size and cost-per-sample without a risk score as follows
N=100000
k1=0.4

# Kernel width and variance for GP
k_width=5000
var_u=8000000

# Suppose we begin with loss estimates at n-values
nset=c(10000,20000,30000)

# with cost-per-individual estimates
k2=c(0.35,0.26,0.28)

# and associated error on those estimates
var_w=c(0.02^2,0.01^2,0.03^2)

# We estimate theta from these three points
theta=powersolve(nset,k2,y_var=var_w)$par

# We will estimate the posterior at these values of n
n=seq(1000,50000,length=1000)

# Mean and variance
p_mu=mu_fn(n,nset=nset,d=k2,var_w = var_w, N=N,k1=k1,theta=theta,
           k_width=k_width,var_u=var_u)
p_var=psi_fn(n,nset=nset,N=N,var_w = var_w,k_width=k_width,var_u=var_u)

# Plot
plot(0,xlim=range(n),ylim=c(20000,60000),type="n",
     xlab="Training/holdout set size",
     ylab="Total cost (= num. cases)")
lines(n,p_mu,col="blue")
```

```

lines(n,p_mu - 3*sqrt(p_var),col="red")
lines(n,p_mu + 3*sqrt(p_var),col="red")
points(nset,k1*nset + k2*(N-nset),pch=16,col="purple")
lines(n,k1*n + powerlaw(n,theta)*(N-n),lty=2)
segments(nset,k1*nset + (k2 - 3*sqrt(var_w))*(N-nset),
         nset,k1*nset + (k2 + 3*sqrt(var_w))*(N-nset))
legend("topright",
      c(expression(mu(n)),
        expression(mu(n) %+-% 3*sqrt(psi(n))),
        "prior(n)",
        "d",
        "3SD(d|n)"),
      lty=c(1,1,2,NA,NA),lwd=c(1,1,1,NA,NA),pch=c(NA,NA,NA,16,124),
      pt.cex=c(NA,NA,NA,1,1),
      col=c("blue","red","black","purple","black"),bg="white")

```

next_n

Finds best value of n to sample next

Description

Recommends a value of n at which to next evaluate individual cost in order to most accurately estimate optimal holdout size. Currently only for use with a power-law parametrisation of k_2 .

Approximately finds a set of n points which, given estimates of cost, minimise width of 95% confidence interval around OHS. Uses a greedy algorithm, so various parameters can be learned along the way.

Given existing training set size/cost estimates $nset$ and d , with $var_w[i]=variance(d[i])$, finds, for each candidate point $n[i]$, the median width of the 90% confidence interval for OHS if

```
nset <-c(nset,n[i]) var_w <-c(var_w,mean(var_w)) d <-c(d,rnorm(powerlaw(n[i],theta),variance=mean(v
```

Usage

```

next_n(
  n,
  nset,
  d,
  N,
  k1,
  nmed = 100,
  var_w = rep(1, length(nset)),
  mode = "asymptotic",
  ...
)

```

Arguments

| | |
|--------|--|
| n | Set of training set sizes to evaluate |
| $nset$ | Training set sizes for which a loss has been evaluated |
| d | Loss at training set sizes $nset$ |
| N | Total number of samples on which the model will be fitted/used |

| | |
|-------|---|
| k1 | Mean loss per sample with no predictive score in place |
| nmed | number of times to re-evaluate d and confidence interval width. |
| var_w | Variance of error in loss estimate at each training set size. |
| mode | Mode for calculating OHS CI (passed to ci_ohs): 'asymptotic' or 'empirical' |
| ... | Passed to powersolve and powersolve_se |

Value

Vector out of same length as n, where out[i] is the expected width of the 95% confidence interval for OHS should n be added to nset.

Examples

```
# Set seed.
set.seed(24015)

# Kernel width and Gaussian process variance
kw0=5000
vu0=1e7

# Include legend on plots or not; inclusion can obscure plot elements on small figures
inc_legend=FALSE

# Suppose we have population size and cost-per-sample without a risk score as follows
N=100000
k1=0.4

# Suppose that true values of a,b,c are given by
theta_true=c(10000,1.2,0.2)
theta_lower=c(1,0.5,0.1) # lower bounds for estimating theta
theta_upper=c(20000,2,0.5) # upper bounds for estimating theta

# We start with five random holdout set sizes (nset0),
# with corresponding cost-per-individual estimates d0 derived
# with various errors var_w0
nstart=10
vwmin=0.001; vwmax=0.005
nset0=round(runif(nstart,1000,N/2))
var_w0=runif(nstart,vwmin,vwmax)
d0=rnorm(nstart,mean=powerlaw(nset0,theta_true),sd=sqrt(var_w0))

# We estimate theta from these three points
theta0=powersolve(nset0,d0,y_var=var_w0,lower=theta_lower,upper=theta_upper,init=theta_true)$par

# We will estimate the posterior at these values of n
n=seq(1000,N,length=1000)

# Mean and variance
p_mu=mu_fn(n,nset=nset0,d=d0,var_w = var_w0, N=N,k1=k1,theta=theta0,k_width=kw0,var_u=vu0)
p_var=psi_fn(n,nset=nset0,N=N,var_w = var_w0,k_width=kw0,var_u=vu0)

# Plot
yrange=c(-30000,100000)
```



```

plot(0,xlim=range(n),ylim=yrange,type="n",
     xlab="Training/holdout set size",
     ylab="Total cost (= num. cases)")
lines(n,p_mu,col="blue")
lines(n,p_mu - 3*sqrt(p_var),col="red")
lines(n,p_mu + 3*sqrt(p_var),col="red")
points(nset0,k1*nset0 + d0*(N-nset0),pch=16,col="purple")
lines(n,k1*n + powerlaw(n,theta0)*(N-n),lty=2)
lines(n,k1*n + powerlaw(n,theta_true)*(N-n),lty=3,lwd=3)
if (inc_legend) {
  legend("topright",
        c(expression(mu(n)),
          expression(mu(n) %+-% 3*sqrt(psi(n))),
          "prior(n)",
          "True",
          "d"),
        lty=c(1,1,2,3,NA),lwd=c(1,1,1,3,NA),pch=c(NA,NA,NA,NA,16),pt.cex=c(NA,NA,NA,NA,1),
        col=c("blue","red","black","purple"),bg="white")
}

## Add line corresponding to recommended new point. This is slow.
nn=seq(1000,N,length=20)
exp_imp <- next_n(nn,nset=nset0,d=d0,var_w = var_w0, N=N,k1=k1,nmed=10,
                  lower=theta_lower,upper=theta_upper)
abline(v=nn[which.min(exp_imp)])

```

ohs_array

*Data for vignette on algorithm comparison***Description**

This object contains data relating to the vignette comparing emulation and parametric algorithms.
 For generation, see hidden code in vignette, or in pipeline at https://github.com/jamesliley/OptHoldoutSize_pipelines

Usage

ohs_array

Format

An object of class array of dimension 200 x 200 x 2 x 2 x 2.

ohs_resample

*Data for vignette on algorithm comparison***Description**

This object contains data relating to the first plot in the vignette comparing emulation and parametric algorithms. For generation, see hidden code in vignette, or in pipeline at https://github.com/jamesliley/OptHoldoutSize_pipelines

Usage

```
ohs_resample
```

Format

An object of class `matrix` (inherits from `array`) with 1000 rows and 4 columns.

`optimal_holdout_size` *Estimate optimal holdout size under parametric assumptions*

Description

Compute optimal holdout size for updating a predictive score given appropriate parameters of cost function

Evaluates empirical minimisation of cost function $l(n; k1, N, \theta) = k1 \cdot n + k2(n; \theta) \cdot (N - n)$.

The function will return `Inf` if no minimum exists. It does not check if the minimum is unique, but this can be guaranteed using the assumptions for theorem 1 in the manuscript.

This calls the function `optimize` from package `stats`.

Usage

```
optimal_holdout_size(N, k1, theta, k2 = powerlaw, round_result = FALSE, ...)
```

Arguments

| | |
|---------------------------|---|
| <code>N</code> | Total number of samples on which the predictive score will be used/fitted. Can be a vector. |
| <code>k1</code> | Cost value in the absence of a predictive score. Can be a vector. |
| <code>theta</code> | Parameters for function <code>k2(n)</code> governing expected cost to an individual sample given a predictive score fitted to <code>n</code> samples. Can be a matrix of dimension <code>n x n_par</code> , where <code>n_par</code> is the number of parameters of <code>k2</code> . |
| <code>k2</code> | Function governing expected cost to an individual sample given a predictive score fitted to <code>n</code> samples. Must take two arguments: <code>n</code> (number of samples) and <code>theta</code> (parameters). Defaults to a power-law form <code>powerlaw(n, c(a, b, c)) = a * n^(-b) + c</code> . |
| <code>round_result</code> | Set to <code>TRUE</code> to solve over integral sizes |
| <code>...</code> | Passed to function <code>optimize</code> |

Value

List/data frame of dimension (number of evaluations) \times (4 + `n_par`) containing input data and results. Columns `size` and `cost` are optimal holdout size and cost at this size respectively. Parameters `N`, `k1`, `theta.1`, `theta.2`, ..., `theta.n_par` are input data.

Examples

```
# Evaluate optimal holdout set size for a range of values of k1 and two values of N, some of which lead to infinity
N1=10000; N2=12000
k1=seq(0.1,0.5,length=20)
A=3; B=1.5; C=0.15; theta=c(A,B,C)

res1=optimal_holdout_size(N1,k1,theta)
res2=optimal_holdout_size(N2,k1,theta)

par(mfrow=c(1,2))
plot(0,type="n",ylim=c(0,500),xlim=range(res1$k1),xlab=expression("k"[1]),ylab="Optimal holdout set size")
  lines(res1$k1,res1$size,col="black")
  lines(res2$k1,res2$size,col="red")
  legend("topright",as.character(c(N1,N2)),title="N:",col=c("black","red"),lty=1)
plot(0,type="n",ylim=c(1500,1600),xlim=range(res1$k1),xlab=expression("k"[1]),ylab="Minimum cost")
  lines(res1$k1,res1$cost,col="black")
  lines(res2$k1,res2$cost,col="red")
  legend("topleft",as.character(c(N1,N2)),title="N:",col=c("black","red"),lty=1)
```

optimal_holdout_size_emulation

Estimate optimal holdout size under semi-parametric assumptions

Description

Compute optimal holdout size for updating a predictive score given a set of training set sizes and estimates of mean cost per sample at those training set sizes.

This is essentially a wrapper for function `mu_fn()`.

Usage

```
optimal_holdout_size_emulation(
  nset,
  d,
  var_w,
  N,
  k1,
  var_u = 1e+07,
  k_width = 5000,
  mean_fn = powerlaw,
  theta = powersolve(nset, d, y_var = var_w)$par,
  npoll = 1000,
  ...
)
```

Arguments

| | |
|--------------------|--|
| <code>nset</code> | Training set sizes for which a loss has been evaluated |
| <code>d</code> | Loss at training set sizes <code>nset</code> |
| <code>var_w</code> | Variance of error in loss estimate at each training set size. |
| <code>N</code> | Total number of samples on which the model will be fitted/used |

| | |
|---------|--|
| k1 | Mean loss per sample with no predictive score in place |
| var_u | Marginal variance for Gaussian process kernel. Defaults to 1e7 |
| k_width | Kernel width for Gaussian process kernel. Defaults to 5000 |
| mean_fn | Functional form governing expected loss per sample given sample size. Should take two parameters: n (sample size) and theta (parameters). Defaults to function powerlaw. |
| theta | Current estimates of parameter values for mean_fn. Defaults to the MLE power-law solution corresponding to n,d, and var_w. |
| npoll | Check npoll equally spaced values between 1 and N for minimum. If NULL, check all values (this can be slow). Defaults to 1000 |
| ... | Passed to function optimise() |

Value

Object of class 'optholdoutsized_emul' with elements "cost" (minimum cost), "size" (OHS), "nset", "d", "var_w", "N", "k1", "theta" (parameters)

Examples

See examples for mu_fn()

| | |
|-------------|---------------------------|
| oracle_pred | <i>Generate responses</i> |
|-------------|---------------------------|

Description

Probably for deprecation

Usage

oracle_pred(X, coefs, num_vars = 3, noise = TRUE)

Arguments

| | |
|----------|--|
| X | Matrix of observations |
| coefs | Vector of coefficients for logistic model. |
| num_vars | If noise==FALSE, computes using only first num_vars predictors |
| noise | If TRUE, uses all predictors |

Value

Vector of predictions

Examples

See examples for model_predict

| | |
|--------------|---|
| params_aspre | <i>Parameters of reported ASPRE dataset</i> |
|--------------|---|

Description

Distribution of covariates for ASPRE dataset; see Rolnik, 2017, NEJM

Usage

```
params_aspre
```

Format

An object of class list of length 16.

| | |
|----------------------|-------------------------------------|
| plot.optholdoutsized | <i>Plot estimated cost function</i> |
|----------------------|-------------------------------------|

Description

Plot estimated cost function, when parametric method is used for estimation.

Draws cost function as a line and indicates minimum. Assumes a power-law form of k2 unless parameter k2 is set otherwise.

Usage

```
## S3 method for class 'optholdoutsized'
plot(x, ..., k2 = powerlaw)
```

Arguments

| | |
|-----|--|
| x | Object of type optholdoutsized |
| ... | Other arguments passed to plot() and lines() |
| k2 | Function governing expected cost to an individual sample given a predictive score fitted to n samples. Must take two arguments: n (number of samples) and theta (parameters). Defaults to a power-law form $\text{powerlaw}(n, c(a, b, c)) = a n^{(-b)} + c$. |

Examples

```
# Simple example

N=100000;
k1=0.3
A=8000; B=1.5; C=0.15; theta=c(A,B,C)

res1=optimal_holdout_size(N,k1,theta)

plot(res1)
```

```
plot.optholdoutsizes_emul
```

Plot estimated cost function using emulation (semiparametric)

Description

Plot estimated cost function, when semiparametric (emulation) method is used for estimation.

Draws posterior mean of cost function as a line and indicates minimum. Also draws mean ± 3 SE.

Assumes a power-law form of k_2 unless parameter k_2 is set otherwise.

Usage

```
## S3 method for class 'optholdoutsizes_emul'
plot(x, ..., k2 = powerlaw)
```

Arguments

| | |
|-----|--|
| x | Object of type optholdoutsizes_emul |
| ... | Other arguments passed to plot() |
| k2 | Function governing expected cost to an individual sample given a predictive score fitted to n samples. Must take two arguments: n (number of samples) and theta (parameters). Defaults to a power-law form $\text{powerlaw}(n, c(a, b, c)) = a n^{-(b) + c}$. |

Examples

```
# Simple example

# Parameters
N=100000;
k1=0.3
A=8000; B=1.5; C=0.15; theta=c(A,B,C)

# True mean function
k2_true=function(n) powerlaw(n,theta)

# Values of n for which cost has been estimated
np=50 # this many points
nset=round(runif(np,1,N))
var_w=runif(np,0.001,0.002)
d=rnorm(np,mean=k2_true(nset),sd=sqrt(var_w))

# Compute OHS
res1=optimal_holdout_size_emulation(nset,d,var_w,N,k1)

# Plot
plot(res1)
```

| | |
|------------|----------------------------|
| powersolve | <i>Fit power law curve</i> |
|------------|----------------------------|

Description

Find least-squares solution: MLE of (a,b,c) under model $y_i = a x_i^{-b} + c + e_i$; $e_i \sim N(0, y_var_i)$

Usage

```
powersolve(
  x,
  y,
  init = c(20000, 2, 0.1),
  y_var = rep(1, length(y)),
  estimate_s = FALSE,
  ...
)
```

Arguments

| | |
|------------|--|
| x | X values |
| y | Y values |
| init | Initial values of (a,b,c) to start. Default c(20000,2,0.1) |
| y_var | Optional parameter giving sampling variance of each y value. Defaults to 1. |
| estimate_s | Parameter specifying whether to also estimate s (as above). Defaults to FALSE (no). |
| ... | further parameters passed to optim. We suggest specifying lower and upper bounds for (a,b,c); e.g. lower=c(1,0,0),upper=c(10000,3,1) |

Value

List (output from optim) containing MLE values of (a,b,c)

Examples

```
# Retrieval of original values
A_true=2000; B_true=1.5; C_true=0.3; sigma=0.002

X=1000*abs(rnorm(10000,mean=4))
Y=A_true*(X^(-B_true)) + C_true + rnorm(length(X),sd=sigma)

c(A_true,B_true,C_true)
powersolve(X[1:10],Y[1:10])$par
powersolve(X[1:100],Y[1:100])$par
powersolve(X[1:1000],Y[1:1000])$par
powersolve(X[1:10000],Y[1:10000])$par
```

| | |
|--------------------|---|
| powersolve_general | <i>General solver for power law curve</i> |
|--------------------|---|

Description

Find least-squares solution: MLE of (a,b,c) under model $y_i = a x_i^{-b} + c + e_i$; $e_i \sim N(0, y_var_i)$

Try a range of starting values and refine estimate.

Slower than a single call to powersolve()

Usage

```
powersolve_general(x, y, y_var = rep(1, length(x)), ...)
```

Arguments

| | |
|-------|--|
| x | X values |
| y | Y values |
| y_var | Optional parameter giving sampling variance of each y value. Defaults to 1. |
| ... | further parameters passed to optim. We suggest specifying lower and upper bounds for (a,b,c); e.g. lower=c(1,0,0),upper=c(10000,3,1) |

Value

List (output from optim) containing MLE values of (a,b,c)

Examples

```
# Retrieval of original values
A_true=2000; B_true=1.5; C_true=0.3; sigma=0.002

X=1000*abs(rnorm(10000,mean=4))
Y=A_true*(X^(-B_true)) + C_true + rnorm(length(X),sd=sigma)

c(A_true,B_true,C_true)
powersolve_general(X[1:10],Y[1:10])$par
powersolve_general(X[1:100],Y[1:100])$par
powersolve_general(X[1:1000],Y[1:1000])$par
powersolve_general(X[1:10000],Y[1:10000])$par
```

| | |
|---------------|--|
| powersolve_se | <i>Standard error matrix for learning curve parameters (power law)</i> |
|---------------|--|

Description

Find approximate standard error matrix for (a,b,c) under power law model for learning curve.

Assumes that

$$y_i = a x_i^{-b} + c + e, \quad e \sim N(0, s^2 y_{\text{var}_i}^2)$$

Standard error can be computed either asymptotically using Fisher information (method='fisher') or bootstrapped (method='bootstrap')

These estimate different quantities: the asymptotic method estimates

$\text{Var}[\text{MLE}(a, b, c) | X, y_{\text{var}}]$

and the bootstrap method estimates

$\text{Var}[\text{MLE}(a, b, c)]$.

Usage

```
powersolve_se(
  x,
  y,
  method = "fisher",
  init = c(20000, 2, 0.1),
  y_var = rep(1, length(y)),
  n_boot = 1000,
  seed = NULL,
  ...
)
```

Arguments

| | |
|--------|--|
| x | X values (typically training set sizes) |
| y | Y values (typically observed cost per individual/sample) |
| method | One of 'fisher' (for asymptotic variance via Fisher Information) or 'bootstrap' (for Bootstrap) |
| init | Initial values of (a,b,c) to start when computing MLE. Default c(20000,2,0.1) |
| y_var | Optional parameter giving sampling variance of each y value. Defaults to 1. |
| n_boot | Number of bootstrap resamples. Only used if method='bootstrap'. Defaults to 1000 |
| seed | Random seed for bootstrap resamples. Defaults to NULL. |
| ... | further parameters passed to optim. We suggest specifying lower and upper bounds; since optim is called on (a*1000 ^{-b} ,b,c), bounds should be relative to this; for instance, lower=c(0,0,0),upper=c(100,3,1) |

Value

Standard error matrix; approximate covariance matrix of $\text{MLE}(a, b, c)$

Examples

```
A_true=10; B_true=1.5; C_true=0.3; sigma=0.1
set.seed(31525)
```

```

X=1+3*rchisq(10000,df=5)
Y=A_true*(X^(-B_true)) + C_true + rnorm(length(X),sd=sigma)

# 'Observations' - 100 samples
obs=sample(length(X),100,rep=FALSE)
Xobs=X[obs]; Yobs=Y[obs]

# True covariance matrix of MLE of a,b,c on these x values
ntest=1000
abc_mat_xfix=matrix(0,ntest,3)
abc_mat_xvar=matrix(0,ntest,3)
E1=A_true*(Xobs^(-B_true)) + C_true
for (i in 1:ntest) {
  Y1=E1 + rnorm(length(Xobs),sd=sigma)
  abc_mat_xfix[i,]=powersolve(Xobs,Y1)$par # Estimate (a,b,c) with same X

  X2=1+3*rchisq(length(Xobs),df=5)
  Y2=A_true*(X2^(-B_true)) + C_true + rnorm(length(Xobs),sd=sigma)
  abc_mat_xvar[i,]=powersolve(X2,Y2)$par # Estimate (a,b,c) with variable X
}

Ve1=var(abc_mat_xfix) # empirical variance of MLE(a,b,c)|X
Vf=powersolve_se(Xobs,Yobs,method='fisher') # estimated SE matrix, asymptotic

Ve2=var(abc_mat_xvar) # empirical variance of MLE(a,b,c)
Vb=powersolve_se(Xobs,Yobs,method='bootstrap') # estimated SE matrix, bootstrap

cat("Empirical variance of MLE(a,b,c)|X\n")
print(Ve1)
cat("\n")
cat("Asymptotic variance of MLE(a,b,c)|X\n")
print(Vf)
cat("\n\n")
cat("Empirical variance of MLE(a,b,c)\n")
print(Ve2)
cat("\n")
cat("Bootstrap-estimated variance of MLE(a,b,c)\n")
print(Vb)
cat("\n\n")

```

psi_fn

Updating function for variance.

Description

Posterior variance for emulator given points n .

Usage

```
psi_fn(n, nset, var_w, N, var_u = 1e+07, k_width = 5000)
```

Arguments

| | |
|----------------------|---|
| <code>n</code> | Set of training set sizes to evaluate at |
| <code>nset</code> | Training set sizes for which a loss has been evaluated |
| <code>var_w</code> | Variance of error in loss estimate at each training set size. |
| <code>N</code> | Total number of samples on which the model will be fitted/used. Only used to rescale <code>var_w</code> |
| <code>var_u</code> | Marginal variance for Gaussian process kernel. Defaults to $1e7$ |
| <code>k_width</code> | Kernel width for Gaussian process kernel. Defaults to 5000 |

Value

Vector Ψ of same length of `n` where $\Psi_i = \text{var}(\text{posterior}(\text{cost}(n_i)))$

Examples

```
# See examples for `mu_fn`
```

| | |
|---------------------|---|
| <code>sens10</code> | <i>Sensitivity at theshold quantile 10%</i> |
|---------------------|---|

Description

Computes sensitivity of a risk score at a threshold at which 10% of samples (or some proportion `pi_int`) are above the threshold.

Usage

```
sens10(Y, Ypred, pi_int = 0.1)
```

Arguments

| | |
|---------------------|--|
| <code>Y</code> | True labels (1 or 0) |
| <code>Ypred</code> | Predictions (univariate; real numbers) |
| <code>pi_int</code> | Compute sensitivity when a proportion <code>pi_int</code> of samples exceed threshold, default 0.1 |

Value

Sensitivity at this threshold

Examples

```
# Simulate
set.seed(32142)

N=1000
X=rnorm(N); Y=rbinom(N,1,prob=logit(X/2))

pi_int=0.1
q10=quantile(X,1-pi_int) # 10% of X values are above this threshold

print(length(which(Y==1 & X>q10))/length(which(X>q10)))
print(sens10(Y,X,pi_int))
```

| | |
|------------------|---|
| sim_random_aspre | <i>Simulate random dataset similar to ASPRE training data</i> |
|------------------|---|

Description

Generate random population of individuals (e.g., newly pregnant women) with given population parameters

Assumes independence of parameter variation. This is not a realistic assumption, but is satisfactory for our purposes.

Usage

```
sim_random_aspre(n, params = NULL)
```

Arguments

| | |
|--------|--------------------|
| n | size of population |
| params | list of parameters |

Value

Matrix of samples

Examples

```
# Load ASPRE related data
data(params_aspre)

X=sim_random_aspre(1000,params_aspre)

print(c(median(X$age),params_aspre$age$median))

print(rbind(table(X$parity)/1000,params_aspre$parity$freq))
```

`split_data`*Split data*

Description

Split data into holdout and intervention sets

Usage

```
split_data(X, frac)
```

Arguments

`X` Matrix of observations

`frac` Fraction of observations to use for the training set

Value

Vector of TRUE/FALSE values (randomised) with proportion `frac` as TRUE

Examples

```
# See examples for model_predict
```

Index

- * **aspre**
 - add_aspre_interactions, 2
 - aspre, 3
 - aspre_k2, 4
 - sens10, 35
 - sim_random_aspre, 36
- * **data,aspre**
 - aspre_emulation, 4
 - aspre_parametric, 5
 - ci_cover_a_yn, 5
 - ci_cover_e_yn, 6
 - data_example_simulation, 10
 - data_nextpoint_em, 10
 - data_nextpoint_par, 11
 - ohs_array, 25
 - ohs_resample, 25
 - params_aspre, 29
- * **emulation**
 - cov_fn, 9
- * **estimation,aspre**
 - powersolve, 31
 - powersolve_general, 32
 - powersolve_se, 32
- * **estimation,emulation**
 - error_ohs_emulation, 11
 - optimal_holdout_size_emulation, 27
- * **estimation**
 - ci_ohs, 6
 - grad_nstar_powerlaw, 17
 - optimal_holdout_size, 26
 - plot.optholdoutsizesize, 29
 - plot.optholdoutsizesize_emul, 30
- * **simulation**
 - gen_base_coefs, 15
 - gen_preds, 16
 - gen_resp, 16
 - model_predict, 19
 - model_train, 20
 - oracle_pred, 28
 - split_data, 37
- add_aspre_interactions, 2
- aspre, 3
- aspre_emulation, 4
- aspre_k2, 4
- aspre_parametric, 5
- ci_cover_a_yn, 5
- ci_cover_e_yn, 6
- ci_ohs, 6
- cov_fn, 9
- data_example_simulation, 10
- data_nextpoint_em, 10
- data_nextpoint_par, 11
- error_ohs_emulation, 11
- exp_imp_fn, 13
- gen_base_coefs, 15
- gen_preds, 16
- gen_resp, 16
- grad_nstar_powerlaw, 17
- logistic, 18
- logit, 18
- model_predict, 19
- model_train, 20
- mu_fn, 21
- next_n, 23
- ohs_array, 25
- ohs_resample, 25
- optimal_holdout_size, 26
- optimal_holdout_size_emulation, 27
- oracle_pred, 28
- params_aspre, 29
- plot.optholdoutsizesize, 29
- plot.optholdoutsizesize_emul, 30
- powerlaw(mu_fn), 21
- powersolve, 31
- powersolve_general, 32
- powersolve_se, 32
- psi_fn, 34
- sens10, 35
- sim_random_aspre, 36
- split_data, 37